



# SAI VOQ System

Israel Meilik,  
June 2020



# SAI VOQ System - Agenda

- Assumptions
- Add route exercise in a VOQ System
- VOQ System pipeline

# Proposal assumptions



# Assumptions

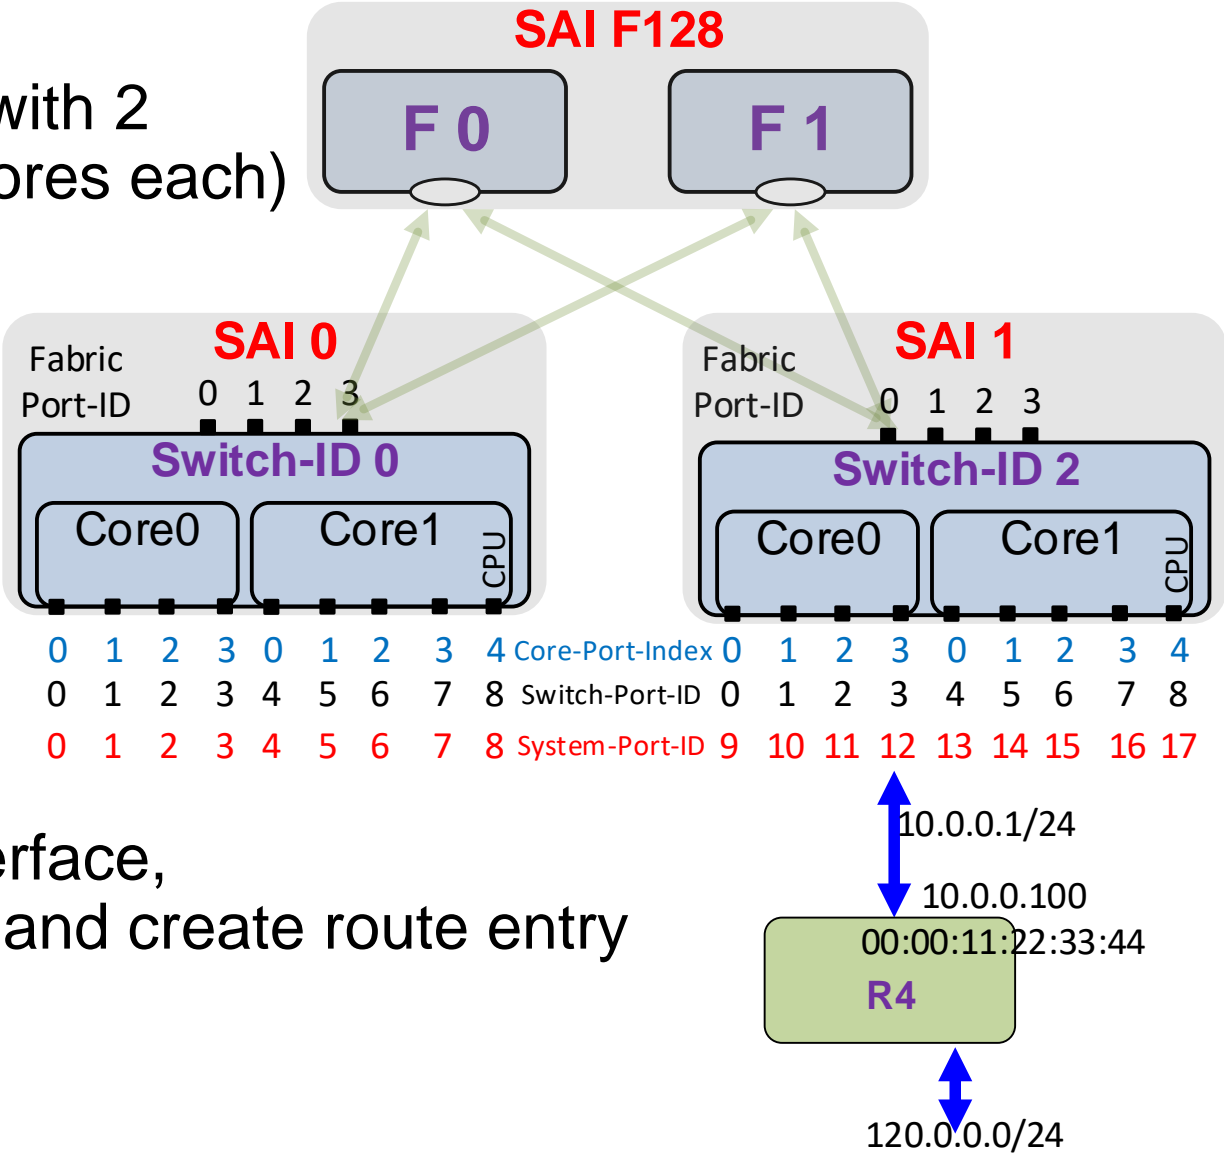
- LAG group has members only on a single VOQ Switch
  - No MLAG
- No IP Multicast
- All Routing Interfaces (RIF) are of type Port, we do not address (Port, VLAN) RIFs
- SAI has an association between System-Port and its local Port (part of SAI VOQ proposal)
- Static System, no ISSU handling/considerations at this phase

# Add Route exercise



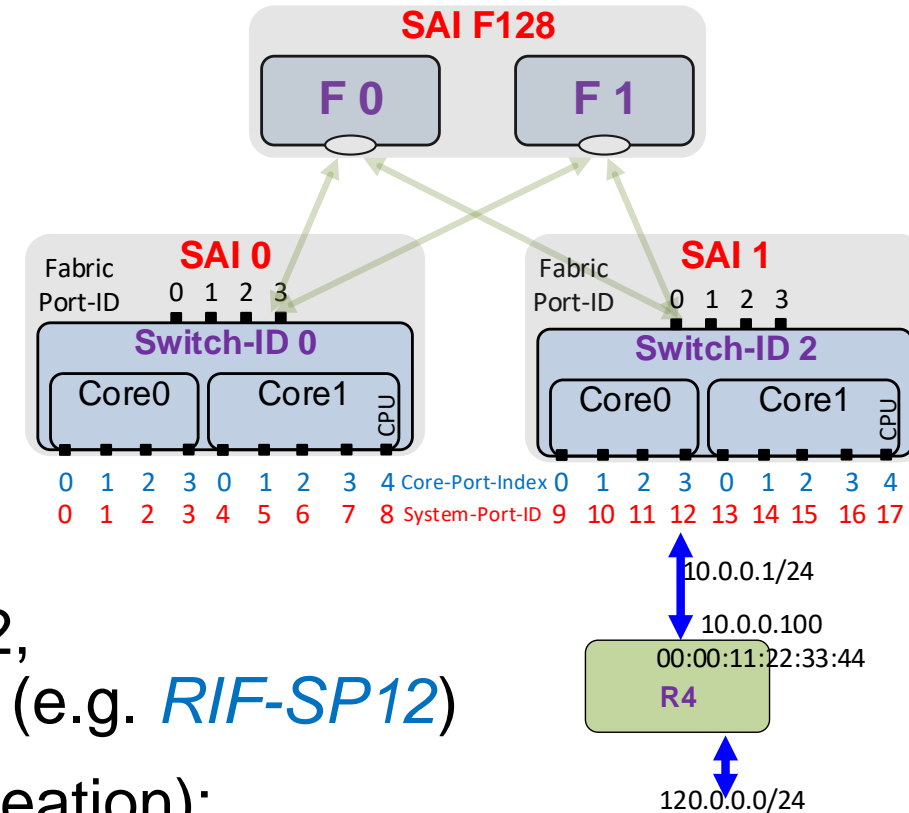
# VOQ Switch System - Add Route exercise

- The exercise will refer to a VOQ system, with 2 Fabric devices, and 2 VOQ Switches (2 cores each)
- Fabric devices are passive (“transparent”) for this packet walk.
- Assumption: upper layer (e.g. SONiC) synchronize information/objects/DBs across all SAI images
- 4 SAI APIs are involved: create router interface, create neighbor entry, create next hop, and create route entry



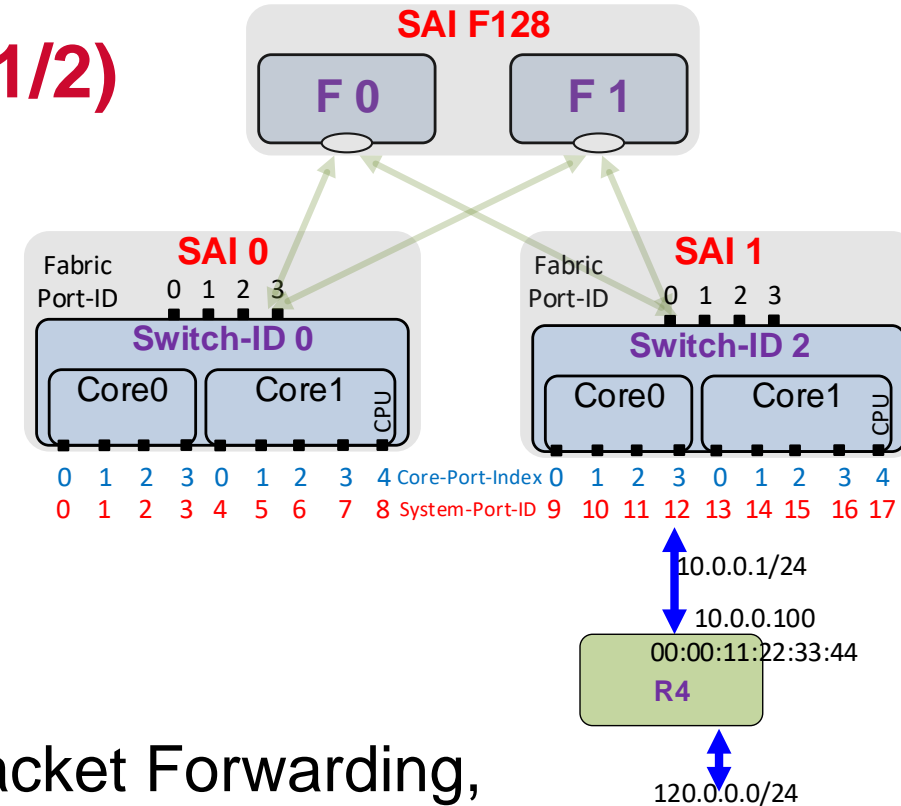
# Add Route - RIF

- Call with **System-Port** instead of Port
- On SAI1:  
*create router Interface (VRF=0, System-Port=12, MAC=00:00:12:34:56:78)*  
⇒ SAI object: *RIF-S1-SP12*
- VOQ-System-DB add a RIF object on system port 12, struct with all the RIF parameters, pointed by handle (e.g. *RIF-SP12*)
- On SAI0 (triggered by upper layer, after *RIF-SP12* creation):  
*create RIF (VRF=0, System-Port=12, MAC=00:00:12:34:56:78)*  
⇒ SAI object: *RIF-S0-SP12*  
Note: function call parameters are provided by upper layer
- In steady state, all RIFs across all system ports exist on every SAI ASIC-DB in the system



## Add Route - Neighbor, functional partition (1/2)

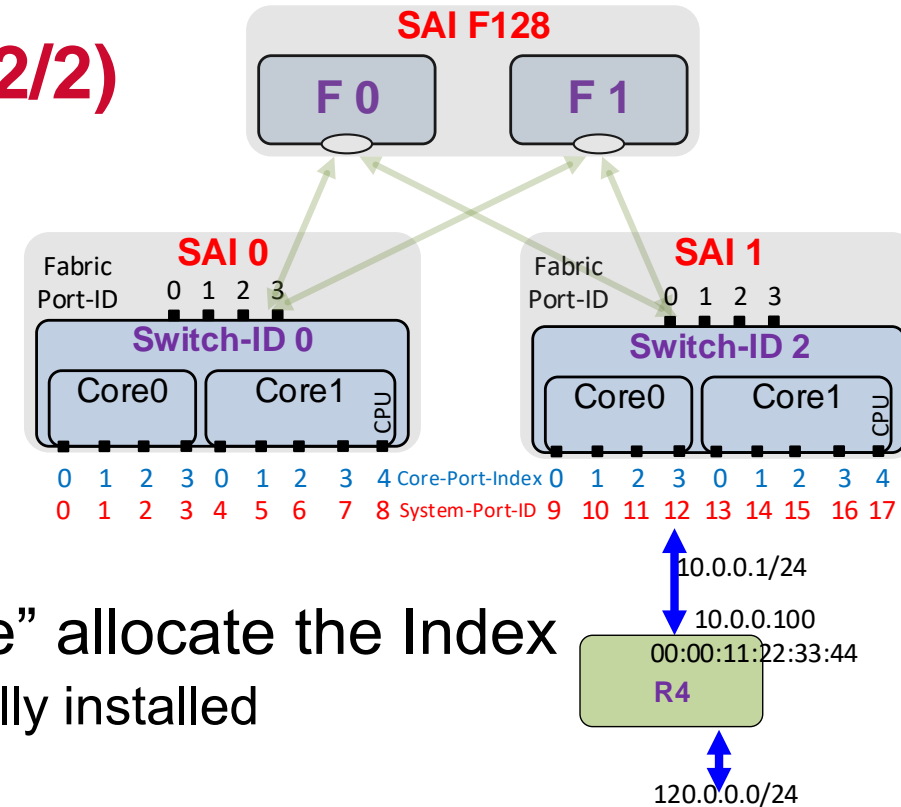
- In a VOQ System, whereas the system behaves as a Single-Hop, in reality the packet is switched through two devices: Ingress Switch and Egress Switch
- This leads to partition of functionality, performed by single Switch in other system, between the Ingress and Egress switches in a VOQ system
- In many architecture the Ingress is responsible for packet Forwarding, i.e. determining the packet destination, and the Egress is responsible for editing the packet
- An alternative architecture of editing the packet on the Ingress device becomes extremely un-efficient in large distributed systems





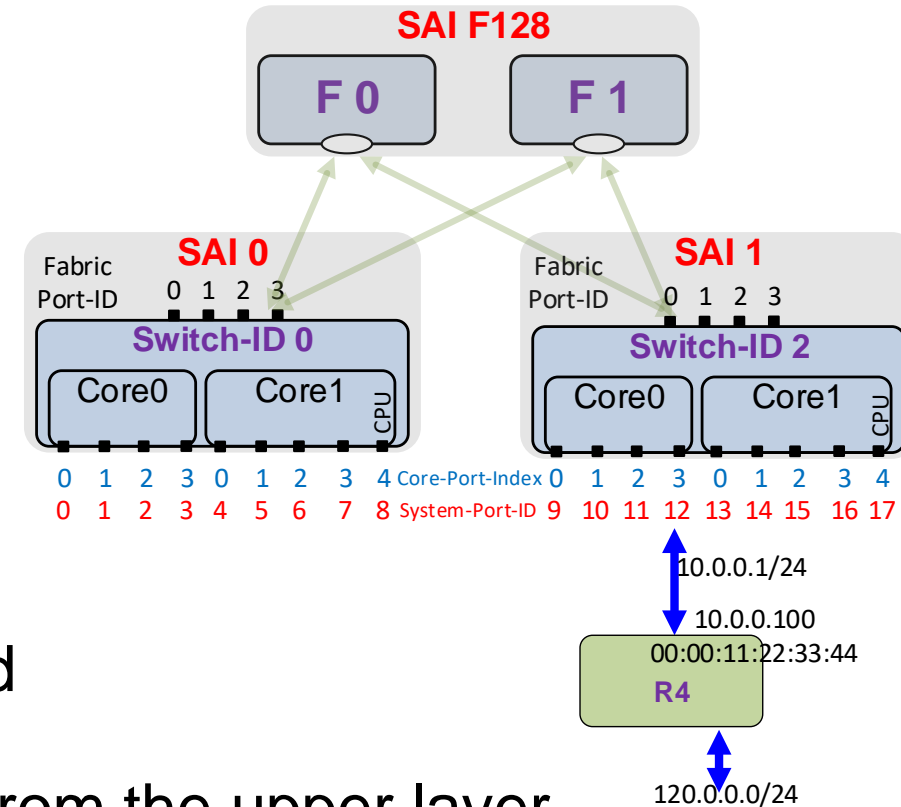
## Add Route - Neighbor, functional partition (2/2)

- To facilitate this functional partition an Index (or few Indexes in other use-cases) is passed from Ingress device to Egress device
- There are few methods for allocation and management of these Indexes
- In some cases it is preferred to let the “owning device” allocate the Index
  - Owner device is the device where the “Egress entry” is actually installed
- In other cases it is preferred to impose the Index (for the “owning device”) by an higher management layer (e.g. SONiC)



# Add Route - Neighbor, new Attributes

- Proposal is to add three new Neighbor attributes:  
SAI-NEIGHBOR-ENCAP-IMPOSE-INDEX // flag  
SAI-NEIGHBOR-PRESENT // flag  
SAI-NEIGHBOR-ENCAP-INDEX // index
- SAI-NEIGHBOR-ENCAP-IMPOSE-INDEX is a flag, determining if the SAI-NEIGHBOR-ENCAP-INDEX will be self allocated (i.e. by the “owning device” SAI), or SAI-NEIGHBOR-ENCAP-INDEX will be imposed from the upper layer
- SAI-NEIGHBOR-PRESENT is a flag, determining if the Neighbor is “installed” on this device, it may be used by some functions such as counters
  - Note that the value of SAI-NEIGHBOR-PRESENT is updated in the event of Neighbor relocation to another Switch



# Add Route - Neighbor

- API Calls:

```
sai_attribute_t remote_attr_list[] = {
```

```
...
```

```
{SAI-NEIGHBOR-ENCAP-INDEX, .value.u32 = index},
```

```
{SAI-NEIGHBOR-PRESENT, .value.bool = true/false}
```

```
{SAI-NEIGHBOR-ENCAP-IMPOSE-INDEX, .value.bool = true/false}
```

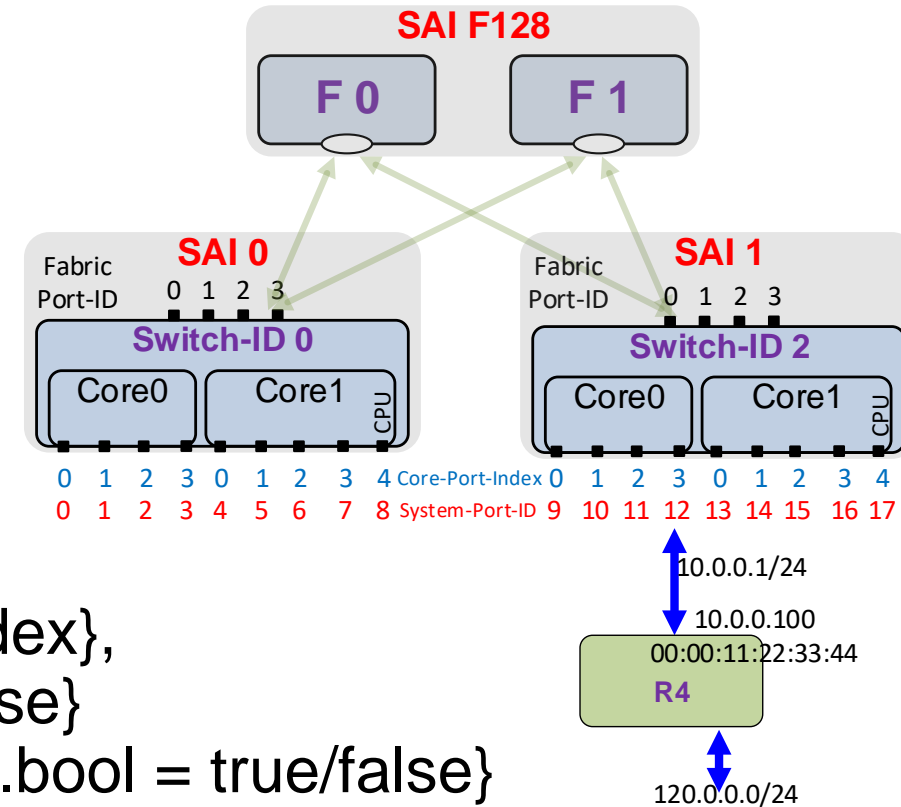
```
...
```

```
};
```

*create neighbor entry (IP-Addr, MAC-Address, RIF,)*

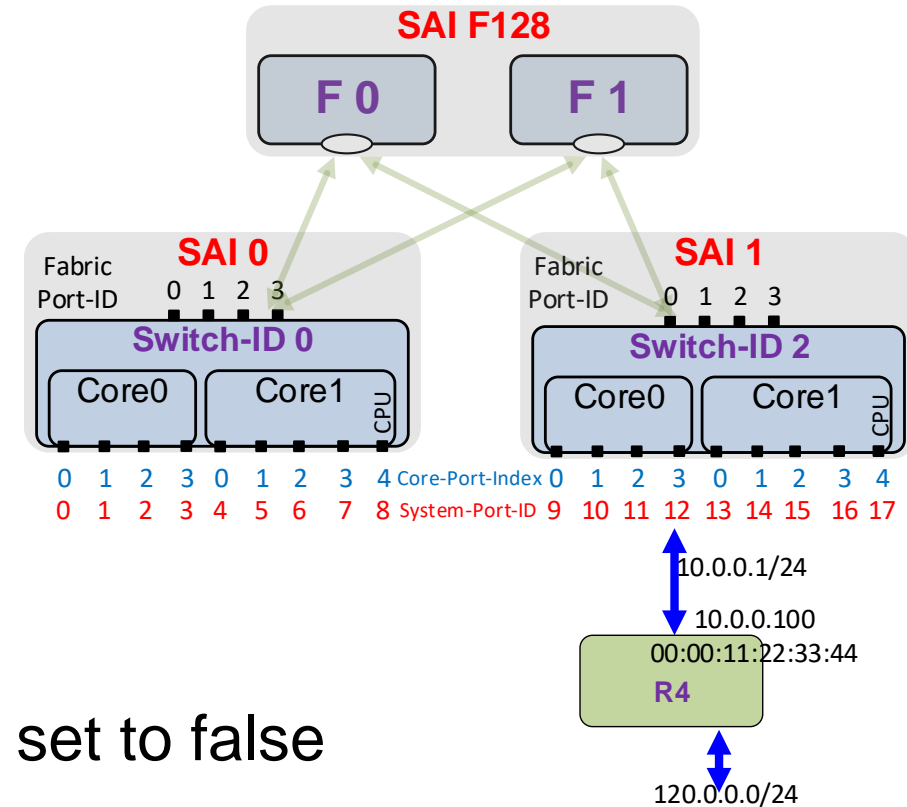
*get neighbor Index (IP-Addr, MAC-Address, RIF)*

- Based on the flag values the SAI will either self allocate or impose the Out-LIF index



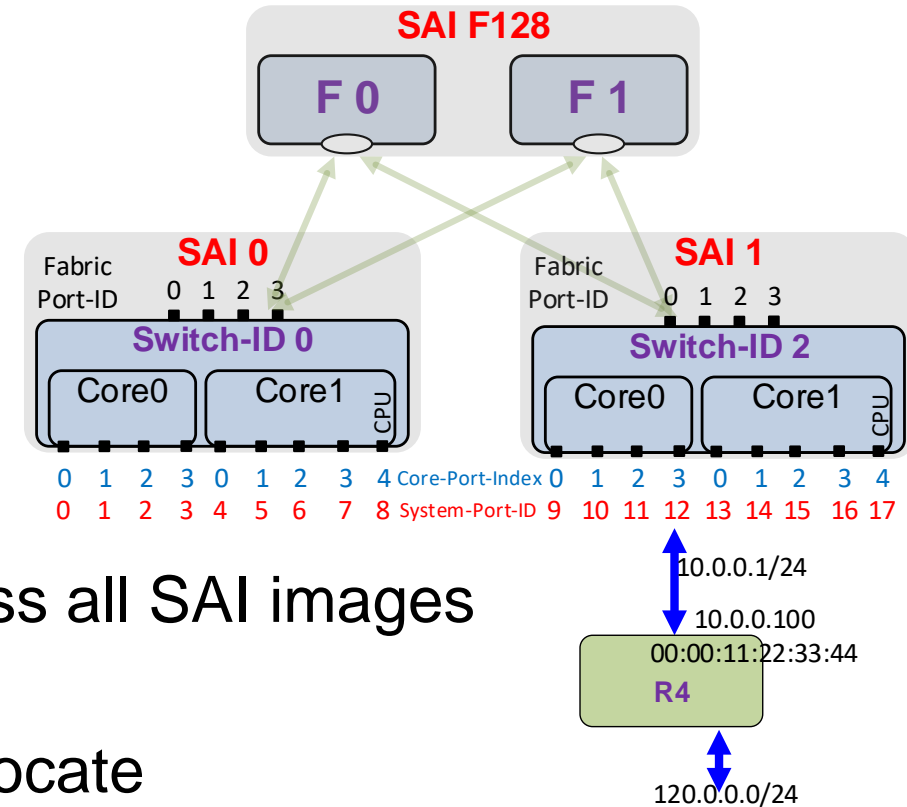
# Add Route - Neighbor, usage method 1

- When the neighbor RIF is port (i.e. exist on a single device), it is recommended that **SAI-NEIGHBOR-ENCAP-INDEX** will be self indexed, i.e. assigned by the owning device SAI
- For example if the RIF on System-Port 12 is a port RIF, then it is expected that SAI of Switch-2 will allocated the **SAI-NEIGHBOR-ENCAP-INDEX**, i.e. **SAI-NEIGHBOR-ENCAP-IMPOSE-INDEX** will be set to false
- Neighbor creation will update the **SAI-NEIGHBOR-ENCAP-INDEX**, which will be synched to other SAIs by upper layer
- Later, when Switch-0 is creating the same neighbor, it is expected that it'll use both **SAI-NEIGHBOR-ENCAP-IMPOSE-INDEX**, and **SAI-NEIGHBOR-ENCAP-INDEX**, in the neighbor creation call



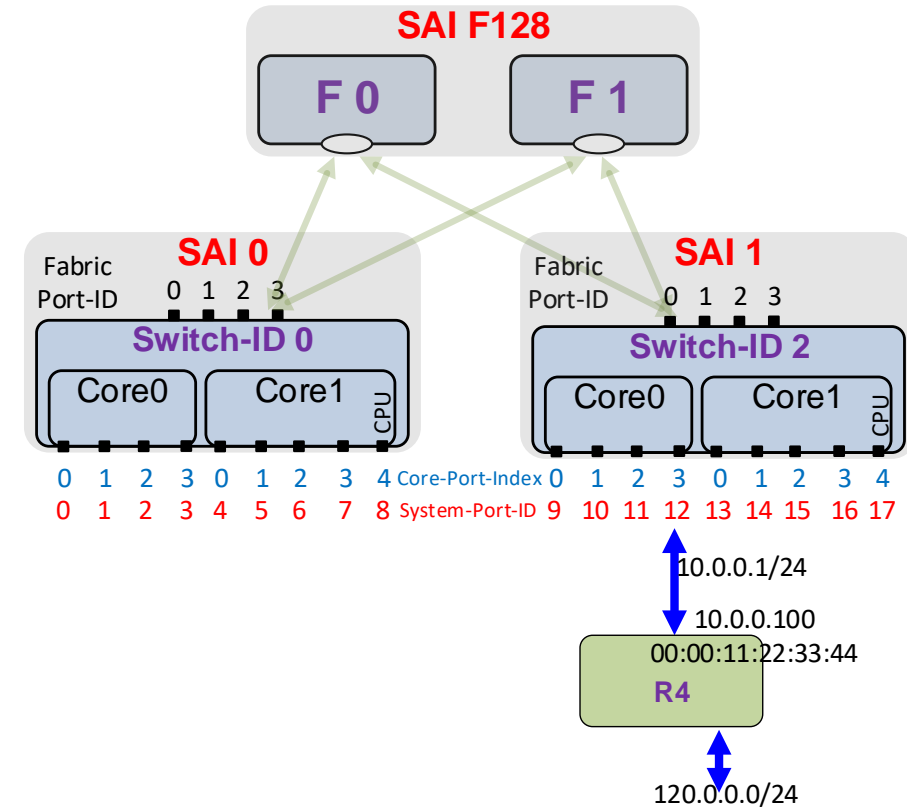
## Add Route - Neighbor, usage method 2

- When the neighbor RIF is VLAN port, or any other case where the neighbor can be relocated to other Switch or Core in the system, it is recommended that **SAI-NEIGHBOR-ENCAP-INDEX** will be imposed by upper layer
- The **SAI-NEIGHBOR-ENCAP-INDEX** is synced across all SAI images
- For example if the RIF on System-Port 12 is a VLAN RIF, then it is expected that upper layer will allocate **SAI-NEIGHBOR-ENCAP-INDEX**, and first impose it on Switch-2, and later on rest of the switches in the system (Switch-0 in this system)



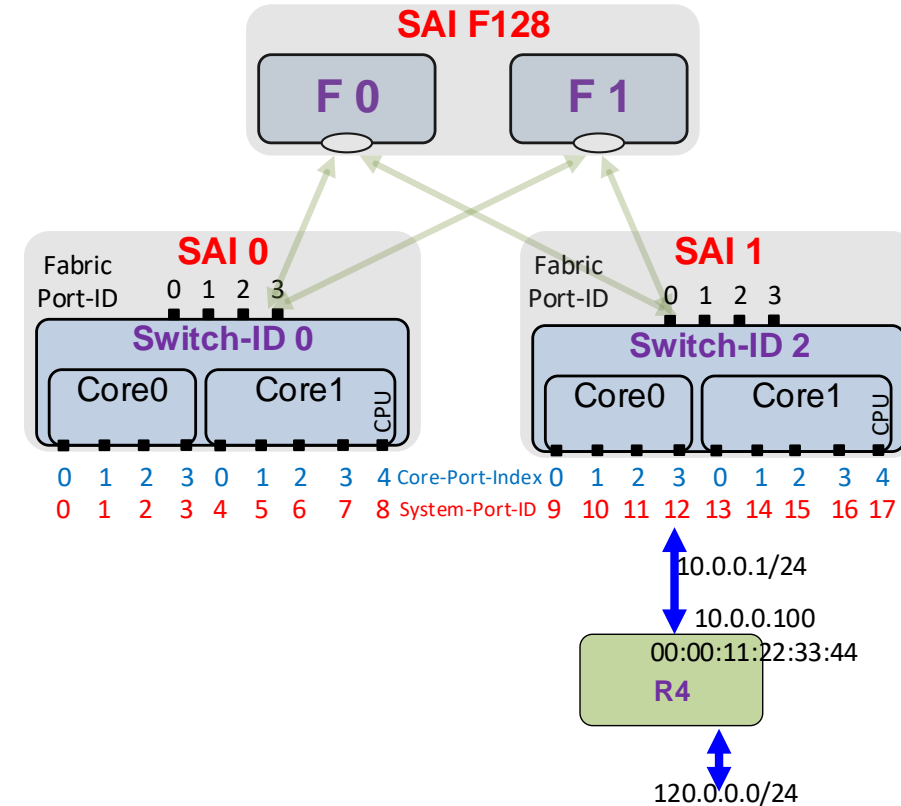
# Add Route - Next Hop

- Next Hop creation is triggered by Neighbor creation
- On SAI1:  
*create next hop (IP=10.0.0.100, RIF=RIF-S1-SP12)*  
⇒ SAI object: *NH-S1-R4*
- On SAI0:  
*create next hop (IP=10.0.0.100, RIF=RIF-S0-SP12)*  
⇒ SAI object: *NH-S0-R4*



# Add Route - Route entry

- Route updates between the different devices in the system is out of scope for this SAI sub group
- On SAI1:  
*create route entry (VRF=0, Prefix=120.0.0.0/24, \*Next-Hop=\*NH-S1-R4)*  
Add prefix to the Router table (LPM) pointing to the FEC associated with *NH-S1-R4*
- On SAI0:  
*create route entry (VRF=0, Prefix=120.0.0.0/24, \*Next-Hop=\*NH-S0-R4)*  
Add prefix to the Router table (LPM) pointing to the FEC associated with *NH-S0-R4*
- No change to create route call



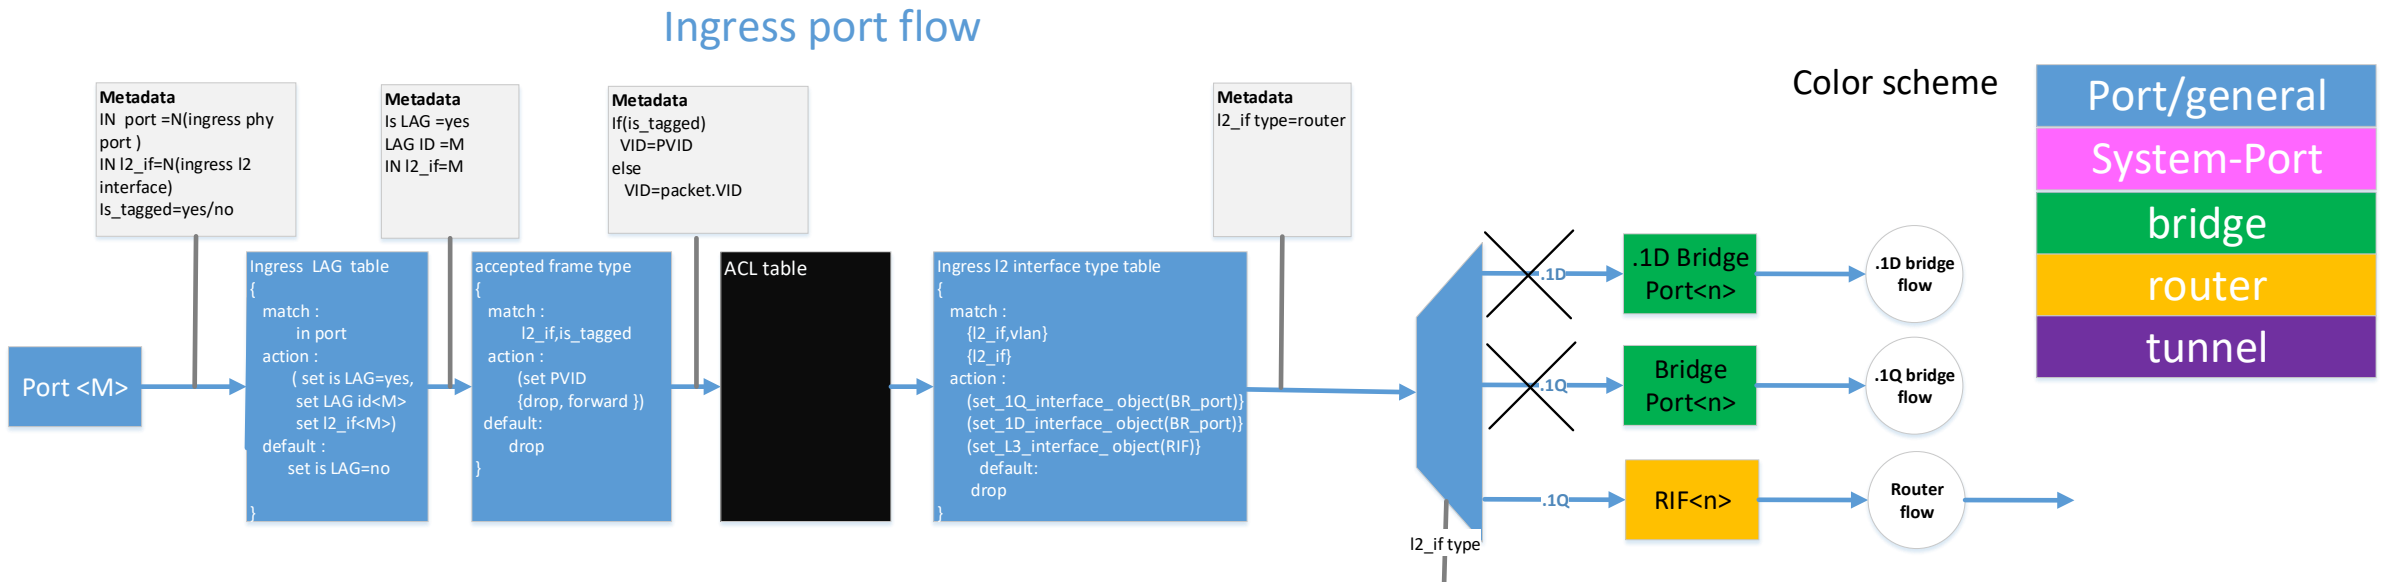
# VOQ System pipeline





# Pipeline - Ingress Port

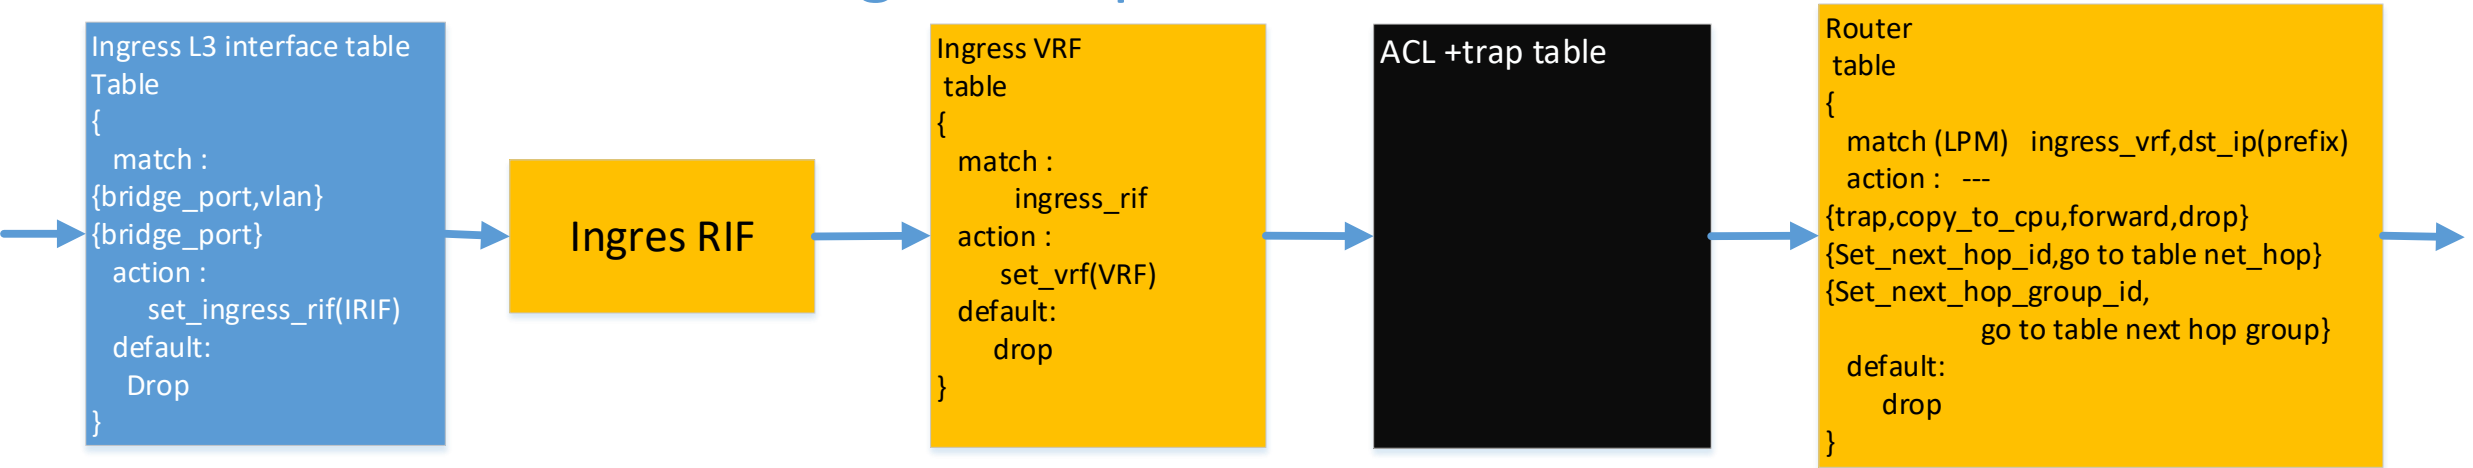
- Split the pipeline into parts, to ease review
- Based on pipeline v7 - UC Routing
- Focus on RIF ports
- No major change



# Pipeline – Route flow, part 1

- No major change

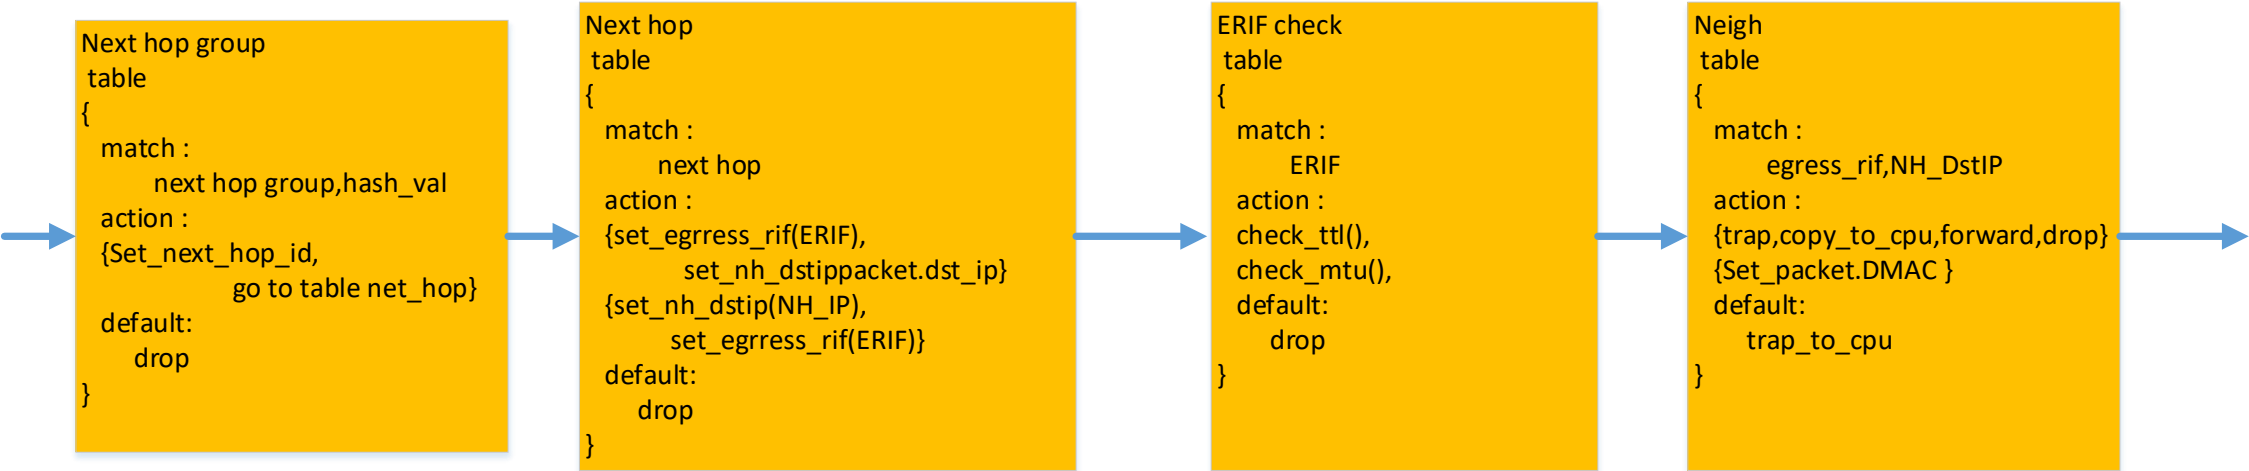
## Routing flow - part 1



# Pipeline - Route flow, part 2 - Existing pipeline

- Note that Neighbor table set the Forwarding decision, but don't set the destination port explicitly

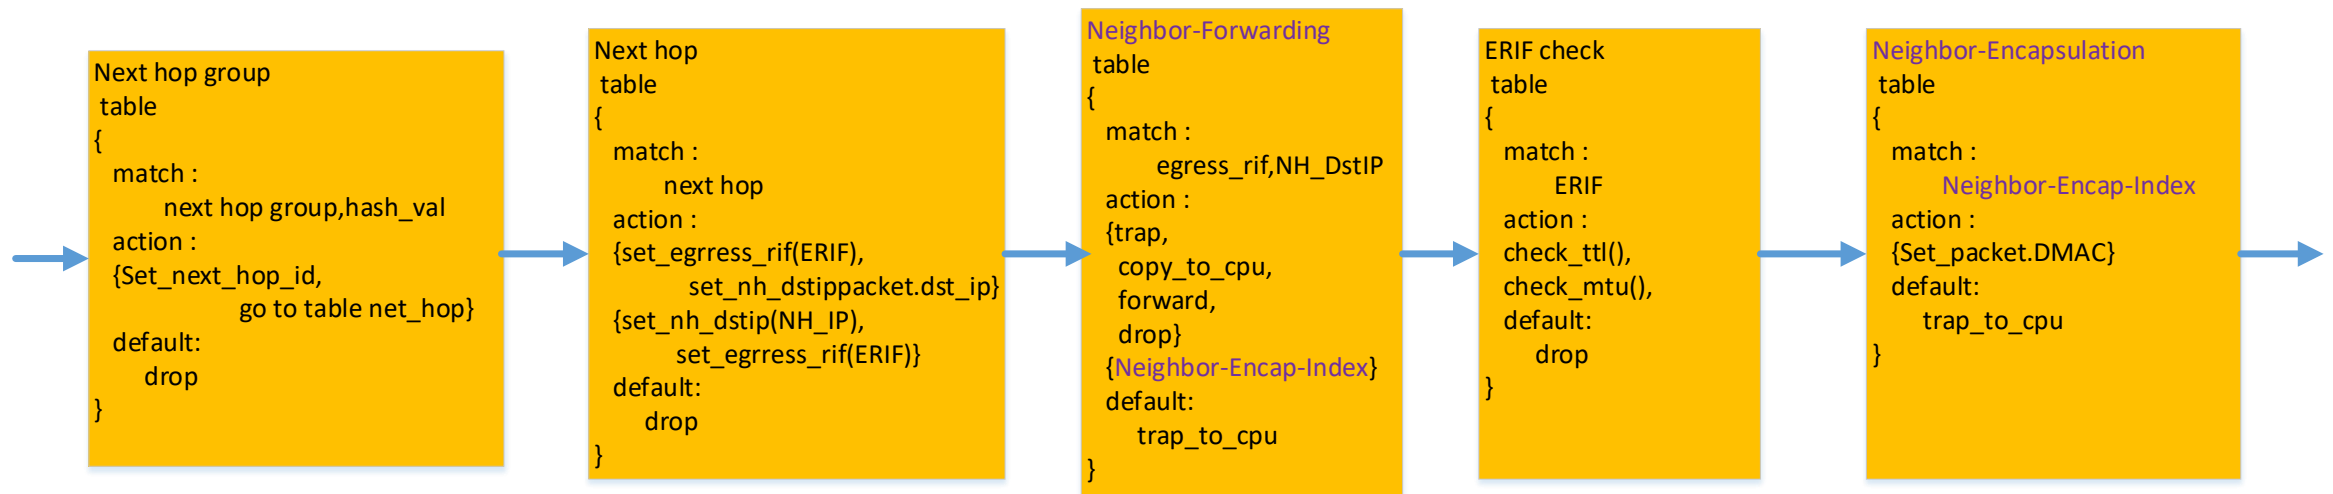
## Routing flow - part 2 (existing)



# Pipeline - Route flow, part 2 - Modified pipeline

- Split Neighbor table into two tables:
  - Neighbor Forwarding
    - Accessed with (egress\_RIF, NH-DstIP) // same as existing
    - decides where to forward the packet: Trap/Snoop/Forward/Drop // same as existing
    - Provides the Index to the Neighbor Encapsulation table
  - Neighbor Encapsulation
    - Provides the packet Dest-MAC // previously in the single Neighbor table

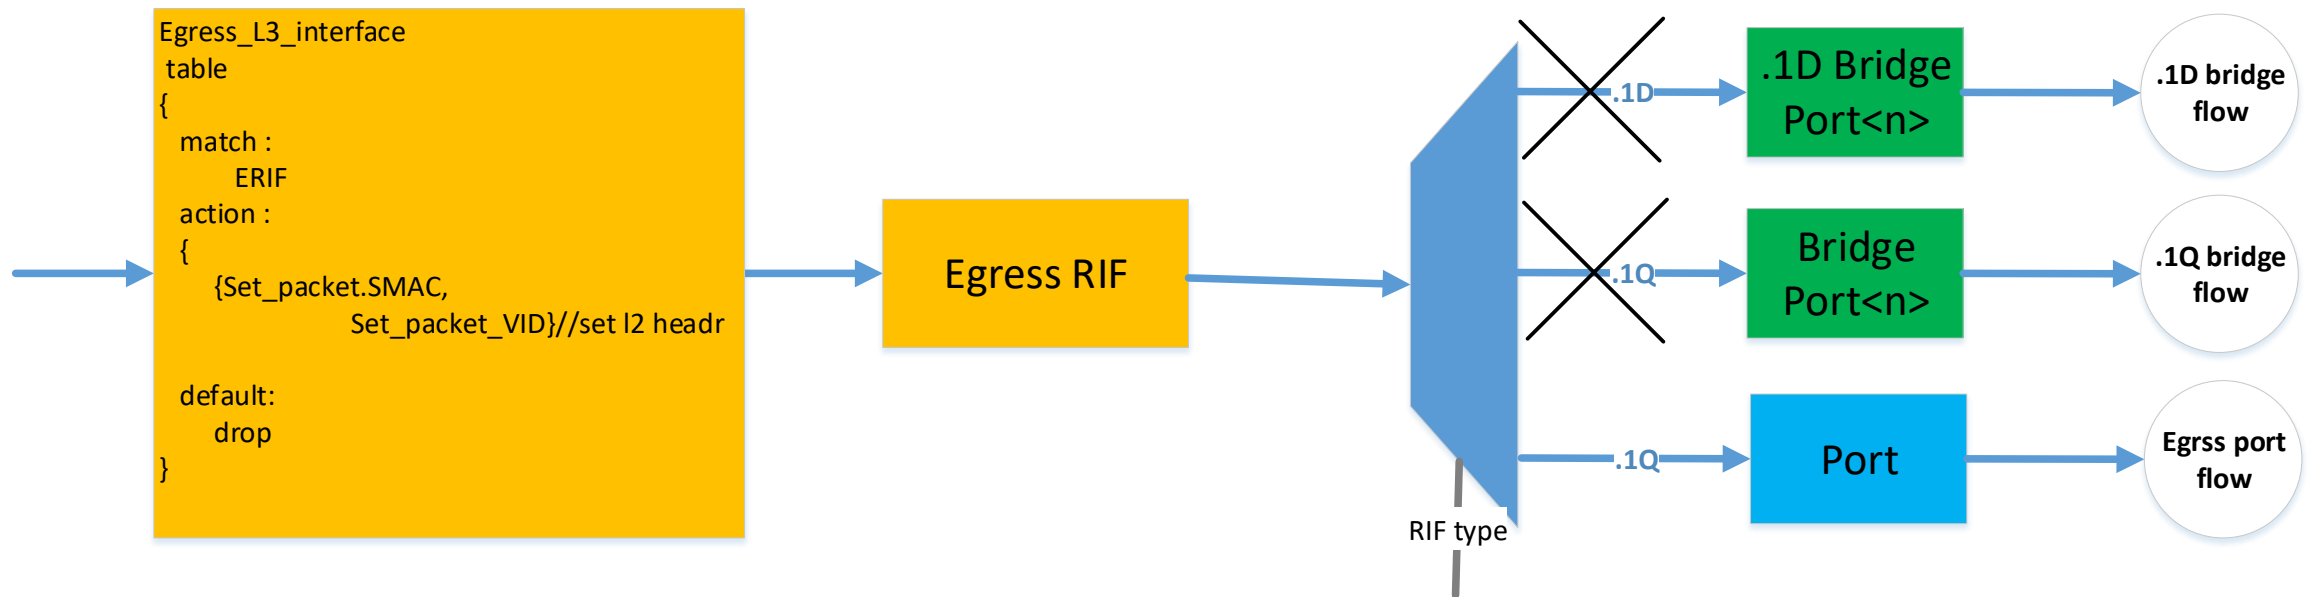
## Routing flow - part 2



# Pipeline – Route flow, part 3

- No major change

## Routing flow - part 3





# Thank You

